

In [1]:

```
# !pip install tensorflow==2.0.0
import tensorflow as tf
```

In [2]:

```
print(tf.__version__) # 2 underlines are required on each side for this to work
```

2.2.0

In [3]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

%matplotlib inline
sns.set(color_codes=True)

from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score
from sklearn.metrics import f1_score, precision_recall_curve, auc

from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, BatchNormalization
from tensorflow.keras import optimizers
```

In [5]:

```
df=pd.read_csv('bank.csv')
```

EDA

General look at head of data, type of content per column, data type per column

1. Check for missing values
2. Remove columns that do not add value to prediction
3. Outlayer treatment

In [6]:

```
display(df.head())
print('The number of rows is', df.shape[0])
print('The number of columns is', df.shape[1])
display(df.count())
print('Are there missing values?', df.isnull().values.any())
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1	15634602	Hargrave	619	France	Female	42	2	
1	2	15647311	Hill	608	Spain	Female	41	1	8380
2	3	15619304	Onio	502	France	Female	42	8	15966
3	4	15701354	Boni	699	France	Female	39	1	
4	5	15737888	Mitchell	850	Spain	Female	43	2	12551

The number of rows is 10000

The number of columns is 14

```
RowNumber      10000
CustomerId      10000
Surname         10000
CreditScore     10000
Geography       10000
Gender          10000
Age            10000
Tenure         10000
Balance        10000
NumOfProducts  10000
HasCrCard      10000
IsActiveMember 10000
EstimatedSalary 10000
Exited         10000
dtype: int64
```

Are there missing values ? False

In [7]:

```
# Another test showing there are no missing values - per column or apply df.isna().sum().sum()
df.count().isnull()
```

Out[7]:

```
RowNumber      False
CustomerId      False
Surname         False
CreditScore     False
Geography       False
Gender          False
Age             False
Tenure          False
Balance         False
NumOfProducts  False
HasCrCard       False
IsActiveMember False
EstimatedSalary False
Exited          False
dtype: bool
```

In [8]:

```
# further exploratory data analysis by checking for unique values in each column
df.nunique()
```

Out[8]:

```
RowNumber      10000
CustomerId      10000
Surname         2932
CreditScore     460
Geography       3
Gender          2
Age             70
Tenure          11
Balance         6382
NumOfProducts   4
HasCrCard       2
IsActiveMember  2
EstimatedSalary 9999
Exited          2
dtype: int64
```

In [9]:

```
# any duplicates to clean data set ? If so apply #df.drop_duplicates(inplace=True)
df.duplicated().any()
```

Out[9]:

```
False
```

In [10]:

```
# check each column to understand its data type before we perform operations
# for example insuring that numbers are not stored as character strings
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   RowNumber             10000 non-null  int64
 1   CustomerId            10000 non-null  int64
 2   Surname               10000 non-null  object
 3   CreditScore           10000 non-null  int64
 4   Geography             10000 non-null  object
 5   Gender               10000 non-null  object
 6   Age                  10000 non-null  int64
 7   Tenure               10000 non-null  int64
 8   Balance              10000 non-null  float64
 9   NumOfProducts        10000 non-null  int64
10   HasCrCard            10000 non-null  int64
11   IsActiveMember       10000 non-null  int64
12   EstimatedSalary      10000 non-null  float64
13   Exited               10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

In [11]:

```
df.head(0)
```

Out[11]:

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
								

In [12]:

```
# DROP COLUMNS WHICH ARE UNIQUE FOR ALL CUSTOMERS AND ADD NO PREDICTIVE VALUE
df.drop('RowNumber', axis=1, inplace=True)
df.drop('CustomerId', axis=1, inplace=True)
df.drop('Surname', axis=1, inplace=True)
```

In [13]:

```
# check if columns have been dropped as expeted
df.head(5)
```

Out[13]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	Is
0	619	France	Female	42	2	0.00	1	1	
1	608	Spain	Female	41	1	83807.86	1	0	
2	502	France	Female	42	8	159660.80	3	1	
3	699	France	Female	39	1	0.00	2	0	
4	850	Spain	Female	43	2	125510.82	1	1	

In [14]:

```
# Examine the data for skewness, outlayers, etc
display(df.describe().T)
```

	count	mean	std	min	25%	50%	
CreditScore	10000.0	650.528800	96.653299	350.00	584.00	652.000	718
Age	10000.0	38.921800	10.487806	18.00	32.00	37.000	44
Tenure	10000.0	5.012800	2.892174	0.00	3.00	5.000	7
Balance	10000.0	76485.889288	62397.405202	0.00	0.00	97198.540	127644
NumOfProducts	10000.0	1.530200	0.581654	1.00	1.00	1.000	2
HasCrCard	10000.0	0.705500	0.455840	0.00	0.00	1.000	1
IsActiveMember	10000.0	0.515100	0.499797	0.00	0.00	1.000	1
EstimatedSalary	10000.0	100090.239881	57510.492818	11.58	51002.11	100193.915	149388
Exited	10000.0	0.203700	0.402769	0.00	0.00	0.000	0

In [15]:

```
# data columns look well distributed. Only EstimatedSalary has a minimum very far from mean and 25% 1st Quartile- an outlayer,
# but we will check for outlayers by visual inspection = boxplots
```

In [16]:

```
# IDENTIFY THE FEATURES AND TARGET VARIABLE
# THE TARGET VARIABLE IS THE ONE THAT IDENTIFIES IF THE CLIENT WILL BE LEAVING THE BANK IN THE NEXT 6 MONTHS !
# HENCE THE TARGET VARIABLE IS THE EXITED COLUMN
# THE REMAINING 8 COLUMNS ARE THE FEATURE COLUMNS
```

In [17]:

```
display(df.columns)
```

```
Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
      'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary',
      'Exited'],
      dtype='object')
```

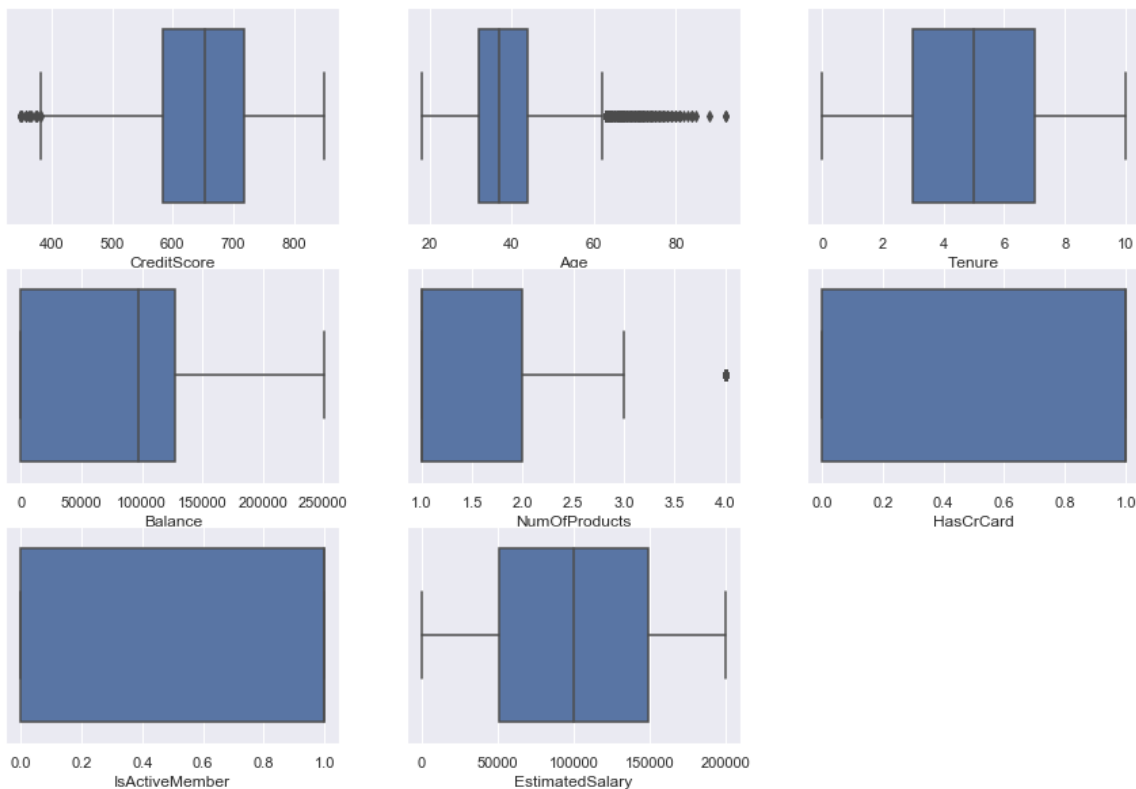
Outlayer Treatment Visualization Check

```
plt.figure(figsize=(15,10)) pos=1 for i in df.columns: plt.subplot(3,3,pos) sns.boxplot(df[i]) pos+=1
```

In [18]:

```
cols=['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary']
# List of independent variables to be checked for outliers

plt.figure(figsize=(15,10))
pos=1
for i in cols:
    plt.subplot(3,3,pos) #---TN: Needs to be updated based on the number of columns. If we have more than 9 columns
    # we need to go to 4x3 grids in the plot aread. Can't have box plot for categorical data
    sns.boxplot(df[i])
    pos+=1
```



In [19]:

```
# treatment of outliers- use the 1.5 IQR
```

In [20]:

```
cols1=['CreditScore','Age', 'Tenure', 'Balance', 'EstimatedSalary']


for i in cols:
    q1,q2,q3=df[i].quantile([0.25,0.5,0.75])
    IQR=q3-q1
    upper_cap=q3+1.5*IQR
    lower_cap=q1-1.5*IQR
    df[i]=df[i].apply(lambda x: upper_cap if x>(upper_cap) else (lower_cap if x<(lower_
cap) else x))
```

In [21]:

```
# now that data has been corrected for outliers perform Univariate Analysis to see if
its still skewed
```

In [22]:

```
plt.figure(figsize=(20,20))
pos=1
for i in cols:
    plt.subplot(3,3,pos)
    sns.distplot(df[i], hist=False)
    pos+=1
```



In [23]:

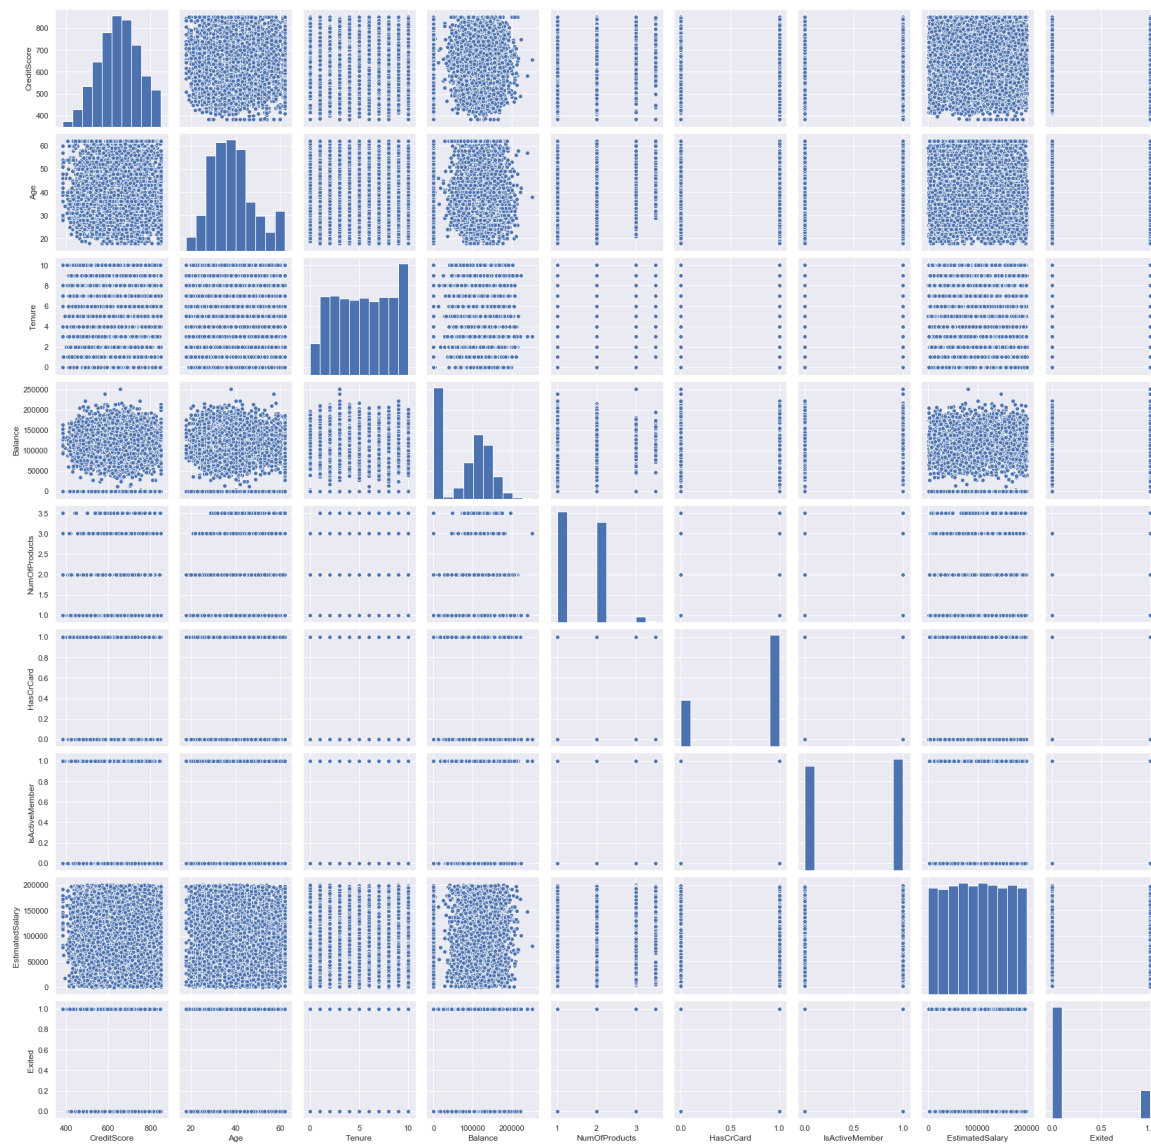
```
# UNIVARIATE ANALYSIS
# Credit score is a bit skewed, so is age- seems baby boomers, probably due to demogra
phis are over 60 and quite a few of them
# Balances, NumOfProducts and HasCrCard are - ideal canidates for clustering, :) but
we are doing supervised Learning here so
# Balances - shows 2 peaks with most people having around zero money (students, unemplo
yed, retired) and quite a few have some
# money a second peak - the working class mid-aged people, which is refileted in Age and
may be correlated with Balances
```

In [24]:

```
# Bivariate Analysis
sns.pairplot(df)
```

Out[24]:

<seaborn.axisgrid.PairGrid at 0x24b9eae8d88>



In [25]:

```
# Bivariate Analysis
# the plots do not show linear correlation, seems Like data displayed is independent fr
om one another
```


In [26]:

```
df.head(1)
```

Out[26]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsA
0	619.0	France	Female	42.0	2	0.0	1.0	1	

In [27]:

```
# drop the dependent variable, we are checking for correlation for independent vars  
df_ind=df.iloc[:,0:10]  
df_ind.head(1)
```

Out[27]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsA
0	619.0	France	Female	42.0	2	0.0	1.0	1	

In [28]:


```
# Check for Correlation in Independent Variables, drop highly correlated ones

corr_data=df_ind.corr().round(2);
fig, ax=plt.subplots()
fig.set_size_inches(12,6)

plt.title ("Heat Map of Correlations")
sns.heatmap(corr_data,annot=True,vmin=-1,vmax=1) # SOMEHOW cmap='YlGnBu' is not working when i add it here
```

Out[28]:

<matplotlib.axes._subplots.AxesSubplot at 0x24ba3074048>



In [29]:

```
# the variables are not correlated, values are far below the 0.5 general threshold - no need to drop any columns
```

In [30]:

df.nunique()

Out[30]:

```
CreditScore      450
Geography         3
Gender           2
Age              45
Tenure           11
Balance          6382
NumOfProducts     4
HasCrCard         2
IsActiveMember    2
EstimatedSalary  9999
Exited            2
dtype: int64
```

Hot Encode Categorical Features

In [31]:

```
#Hot Encode Categorical columns
df=pd.get_dummies(df, drop_first=True)
```

JUST NOTES FOR MYSELF - PLEASE IGNORE
 df=pd.get_dummies(df['Geography'], columns=['Geography'])
 df=pd.get_dummies(df, columns=['Gender'])

```
data_dummies=pd.get_dummies(df['Geography'],prefix=['Country'], columns=['Geography'])
data_dummies1=pd.get_dummies(df['Gender'],prefix=['Gender'], columns=['Gender'])
print(data_dummies.head(2)) data_dummies1.head(2)
```

df=pd.get_dummies(df) # why was the gender column also not hot encoded since it was an object column ?
 does this not need to go back to main dataframe with the original column being replaced (dropped) by
 dummies ? do i need to specifically hot encode for gender like so - df=pd.get_dummies(df, columns=
 ['Gender']) (which is not working)

In [32]:

```
# visual view of the DF with dummies
df.tail()
```

Out[32]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	Es
9995	771.0	39.0	5	0.00	2.0	1	0	
9996	516.0	35.0	10	57369.61	1.0	1	1	
9997	709.0	36.0	7	0.00	1.0	0	1	
9998	772.0	42.0	3	75075.31	2.0	1	0	
9999	792.0	28.0	4	130142.79	1.0	1	0	

Divide the Data Set into Training and Test Set

In [33]:

```
# dont we need to have a validation set of data ? should we not split data into 3 rather than 2. Why only 20% not 30% below ?
# do all input variables must be float32 or same data type ? How do we change all numerical columns to one datatype ?
X=df.drop(['Exited'], axis=1) # set of independent variables
Y=df['Exited']

X_train,X_test, y_train,y_test=train_test_split(X,Y, test_size=0.20, random_state=1)
```

In [34]:

```
# I HAVE SET MY DEPENDENT VARIABLE AS EXITED. I SEE IN MENTOR SESSION FRAUD DETECTION Y  
HAS BEEN DEFINED AS  
# y_data = credit_data.iloc[:, -1] - IS THIS JUST AN EMPTY SET WITH NO ROWS ? MUST I D  
O THAT ?
```

In [35]:

```
# we can see splitting of data has worked well- 8000 are in training set  
X_train.shape
```

Out[35]:

```
(8000, 11)
```

In [36]:

```
# Scale (Normalize) the data    # why dont we scale the y ? it is a supervised learning  
technique so we have to compare  
# the results of actual y with predicted and must be on same scale ...  
## WE DO NOT SCALE THE y_train or y_test since we are looking to compare predicted 0/1  
with actual (unscaled) 0/1.  
# Besides how wouldl one scale 0 and 1-ones ?  
  
scaler=preprocessing.MinMaxScaler()  
  
X_train=scaler.fit_transform(X_train)  
X_test=scaler.fit_transform(X_test)
```

In [37]:

```
X_train.shape
```

Out[37]:

```
(8000, 11)
```

In [38]:

```
# Initiate construction of ANN
```

In [39]:

```
#model= Sequential()
```

In [41]:

```
# add Layers
model = Sequential(layers=None, name=None) ##-TM: Step 1: Empty Model

model.add(Dense(8,input_shape = (11,), activation = 'relu')) ##-TM: Step 2: Hidden Layer 1; First input is number of neuron
##-TM: Step 2, Second input is number of input variables, i.e. 11; Refer to X_train.shape
model.add(Dense(4,activation='relu'))
model.add(Dense(1,activation='sigmoid')) # Output Layer

sgd=optimizers.Adam(lr=0.001)

model.compile (loss='binary_crossentropy', optimizer=sgd, metrics=['accuracy'])
```

In []:

```
# do we need this ?
# tf.keras.initializers.random_normal
# sgd=optimizers.Adam(lr=0.001)
```

In []:

```
# Training the Model
```

In [42]:

```
model.fit(X_train,y_train.values, epochs=30, batch_size=2000, verbose=1)
# verbose shows output, epochs is number of runs to reduce error
```

```
Epoch 1/30
4/4 [=====] - 0s 3ms/step - loss: 0.5964 - accuracy: 0.7794
Epoch 2/30
4/4 [=====] - 0s 2ms/step - loss: 0.5881 - accuracy: 0.7864
Epoch 3/30
4/4 [=====] - 0s 0s/step - loss: 0.5802 - accuracy: 0.7891
Epoch 4/30
4/4 [=====] - 0s 1ms/step - loss: 0.5725 - accuracy: 0.7918
Epoch 5/30
4/4 [=====] - 0s 1ms/step - loss: 0.5653 - accuracy: 0.7945
Epoch 6/30
4/4 [=====] - 0s 1ms/step - loss: 0.5583 - accuracy: 0.7954
Epoch 7/30
4/4 [=====] - 0s 2ms/step - loss: 0.5517 - accuracy: 0.7956
Epoch 8/30
4/4 [=====] - 0s 2ms/step - loss: 0.5453 - accuracy: 0.7961
Epoch 9/30
4/4 [=====] - 0s 1ms/step - loss: 0.5392 - accuracy: 0.7966
Epoch 10/30
4/4 [=====] - 0s 3ms/step - loss: 0.5335 - accuracy: 0.7971
Epoch 11/30
4/4 [=====] - 0s 1ms/step - loss: 0.5281 - accuracy: 0.7971
Epoch 12/30
4/4 [=====] - 0s 2ms/step - loss: 0.5230 - accuracy: 0.7972
Epoch 13/30
4/4 [=====] - 0s 1ms/step - loss: 0.5184 - accuracy: 0.7975
Epoch 14/30
4/4 [=====] - 0s 2ms/step - loss: 0.5141 - accuracy: 0.7974
Epoch 15/30
4/4 [=====] - 0s 2ms/step - loss: 0.5103 - accuracy: 0.7974
Epoch 16/30
4/4 [=====] - 0s 2ms/step - loss: 0.5069 - accuracy: 0.7974
Epoch 17/30
4/4 [=====] - 0s 1ms/step - loss: 0.5038 - accuracy: 0.7974
Epoch 18/30
4/4 [=====] - 0s 2ms/step - loss: 0.5011 - accuracy: 0.7974
Epoch 19/30
4/4 [=====] - 0s 2ms/step - loss: 0.4987 - accuracy: 0.7974
Epoch 20/30
4/4 [=====] - 0s 2ms/step - loss: 0.4966 - accuracy: 0.7974
Epoch 21/30
```

```
4/4 [=====] - 0s 2ms/step - loss: 0.4947 - accuracy: 0.7974
Epoch 22/30
4/4 [=====] - 0s 1ms/step - loss: 0.4929 - accuracy: 0.7975
Epoch 23/30
4/4 [=====] - 0s 1ms/step - loss: 0.4913 - accuracy: 0.7975
Epoch 24/30
4/4 [=====] - 0s 1ms/step - loss: 0.4898 - accuracy: 0.7975
Epoch 25/30
4/4 [=====] - 0s 3ms/step - loss: 0.4884 - accuracy: 0.7977
Epoch 26/30
4/4 [=====] - 0s 1ms/step - loss: 0.4869 - accuracy: 0.7980
Epoch 27/30
4/4 [=====] - 0s 985us/step - loss: 0.4855 - accuracy: 0.7984
Epoch 28/30
4/4 [=====] - 0s 2ms/step - loss: 0.4842 - accuracy: 0.7993
Epoch 29/30
4/4 [=====] - 0s 3ms/step - loss: 0.4829 - accuracy: 0.7999
Epoch 30/30
4/4 [=====] - 0s 3ms/step - loss: 0.4817 - accuracy: 0.8006
```

Out[42]:

<tensorflow.python.keras.callbacks.History at 0x24ba3004e08>

In []:

```
# Evaluation of the model
```

In [45]:

```
results=model.evaluate(X_test,y_test.values)
```

```
63/63 [=====] - 0s 1ms/step - loss: 0.4879 - accuracy: 0.7925
```

In [49]:

```
print(model.metrics_names)
print(results)
```

```
['loss', 'accuracy']
[0.4879456162452698, 0.7925000190734863]
```


In [51]:

```
#--Creating a data frame to store model results
resultsDf = pd.DataFrame({'Iteration': None, 'Method':['Adam - Check '], 'Log Loss': results[0], 'Accuracy': results[1]})
resultsDf = resultsDf[['Iteration', 'Method', 'Log Loss', 'Accuracy']]
resultsDf
```

Out[51]:

	Iteration	Method	Log Loss	Accuracy
0	None	Adam - Check	0.487946	0.7925

In [52]:

```
Y_pred_cls = model.predict_classes(X_test, batch_size=2000, verbose=0)

print('Accuracy Model1 (Dropout): ' + str(model.evaluate(X_test,y_test.values)[1]))

print('Recall_score: ' + str(recall_score(y_test.values,Y_pred_cls)))

print('Precision_score: ' + str(precision_score(y_test.values, Y_pred_cls)))

print('F-score: ' + str(f1_score(y_test.values,Y_pred_cls)))
confusion_matrix(y_test.values, Y_pred_cls)
```

WARNING:tensorflow:From <ipython-input-52-b98c2c46a761>:1: Sequential.predict_classes (from tensorflow.python.keras.engine.sequential) is deprecated and will be removed after 2021-01-01.

Instructions for updating:

Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if your model does multi-class classification (e.g. if it uses a `softmax` last-layer activation). * `(model.predict(x) > 0.5).astype("int32")`, if your model does binary classification (e.g. if it uses a `sigmoid` last-layer activation).

63/63 [=====] - 0s 1ms/step - loss: 0.4879 - accuracy: 0.7925

Accuracy Model1 (Dropout): 0.7925000190734863

Recall_score: 0.007228915662650603

Precision_score: 0.5

F-score: 0.014251781472684086

Out[52]:

```
array([[1582, 3],
       [ 412, 3]], dtype=int64)
```

In []:

```
# HYPERTUNING
# Experiment with number of neurons, & number of layers, optimizers and activation function to see if you can improve results
```

In [56]:

```
# sigmoid is placed here as the output node function.
def create_model(opt):

    # create model
    model = Sequential(layers=None, name=None)
    model.add(Dense(8, input_shape = (11,), activation = 'relu'))#we are using a different activation function here than above
    model.add(Dense(4, activation = 'relu')) # create second hidden layer
    model.add(Dense(1, activation = 'sigmoid')) # single binary output node

    # Compile model
    model.compile(optimizer = opt, loss = 'binary_crossentropy', metrics=['accuracy'])

    return model
```

In [59]:

```
# try it with different optimizers to Hypertune completely
optimizer=['Adam','SGD','RMSProp','Adadelta','Adagrad','Adamax']
```

In [60]:

```
# 11 PICKED AS THE OPTIMAL NUMBER OF ITERATIONS IS THE INPUT VARIABLE
for j in range(1,11):

    print("Iteration:",j)

    for i in optimizer:

        mod=create_model(i)
        #print("Model:",i,end=' ')
        mod.fit(X_train, y_train.values, batch_size = 700, epochs = j, verbose = 0)
        results = mod.evaluate(X_test, y_test.values, verbose=0)
        #print(mod.metrics_names)
        #print(results)

        tempResultsDf = pd.DataFrame({'Iteration': j, 'Method':[i], 'Accuracy': results[1], 'Log Loss': results[0]})
        resultsDf = pd.concat([resultsDf, tempResultsDf])
        resultsDf = resultsDf[['Iteration', 'Method', 'Accuracy', 'Log Loss']]
        #resultsDf
```

```
Iteration: 1
Iteration: 2
Iteration: 3
Iteration: 4
Iteration: 5
Iteration: 6
Iteration: 7
Iteration: 8
Iteration: 9
Iteration: 10
```

In [61]:

```
resultsDf=resultsDf[resultsDf['Iteration']>0]
resultsDf
```

Out[61]:

	Iteration	Method	Accuracy	Log Loss
0	1	Adadelta	0.7925	0.625028
0	1	Adam	0.7715	0.592211
0	1	SGD	0.6260	0.695954
0	1	RMSProp	0.7910	0.654466
0	1	Adadelta	0.7360	0.656038
...
0	10	SGD	0.7880	0.616201
0	10	RMSProp	0.7925	0.481865
0	10	Adadelta	0.3900	0.762319
0	10	Adagrad	0.7925	0.627326
0	10	Adamax	0.7920	0.613519

61 rows × 4 columns

In [62]:

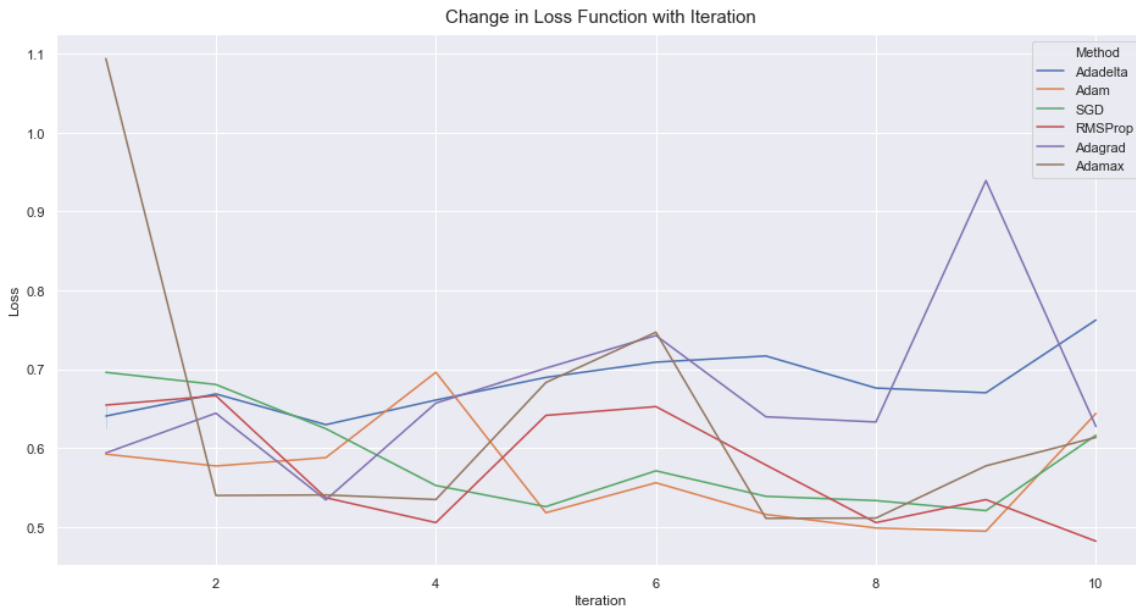
```
# A visual of the results of error reduction and approaching a minima

fig, ax = plt.subplots(figsize=(16,8)); #-----Setting size of the canvas

plt.title('Change in Loss Function with Iteration', fontsize=15, pad=10)
#-----Title of the chart

sns.lineplot(x=resultsDf['Iteration'],y=resultsDf['Log Loss'],hue=resultsDf['Method']);

ax.set(xlabel='Iteration', ylabel='Loss');
```



In [64]:

```
resultsDf_temp=resultsDf[resultsDf['Method']!='Adadelta']  
resultsDf_temp=resultsDf_temp[resultsDf_temp['Method']!='Adagrad']  
resultsDf_temp
```

Out[64]:

	Iteration	Method	Accuracy	Log Loss
0	1	Adam	0.7715	0.592211
0	1	SGD	0.6260	0.695954
0	1	RMSProp	0.7910	0.654466
0	1	Adamax	0.2075	1.093744
0	2	Adam	0.7925	0.577205
0	2	SGD	0.5935	0.680502
0	2	RMSProp	0.7625	0.666256
0	2	Adamax	0.7925	0.539828
0	3	Adam	0.7925	0.587831
0	3	SGD	0.7650	0.624506
0	3	RMSProp	0.7925	0.536941
0	3	Adamax	0.7925	0.540527
0	4	Adam	0.6480	0.696008
0	4	SGD	0.7925	0.552441
0	4	RMSProp	0.7925	0.505408
0	4	Adamax	0.7925	0.534682
0	5	Adam	0.7925	0.517907
0	5	SGD	0.7925	0.525671
0	5	RMSProp	0.7900	0.641346
0	5	Adamax	0.6825	0.683102
0	6	Adam	0.7925	0.555918
0	6	SGD	0.7925	0.571076
0	6	RMSProp	0.7925	0.652504
0	6	Adamax	0.3755	0.746795
0	7	Adam	0.7925	0.515731
0	7	SGD	0.7925	0.538752
0	7	RMSProp	0.7925	0.578463
0	7	Adamax	0.7925	0.510828
0	8	Adam	0.7925	0.498588
0	8	SGD	0.7925	0.533340
0	8	RMSProp	0.7925	0.505403
0	8	Adamax	0.7925	0.511161
0	9	Adam	0.7925	0.494567
0	9	SGD	0.7920	0.520574
0	9	RMSProp	0.7925	0.534559
0	9	Adamax	0.7925	0.577384
0	10	Adam	0.7585	0.643810

	Iteration	Method	Accuracy	Log Loss
0	10	SGD	0.7880	0.616201
0	10	RMSProp	0.7925	0.481865
0	10	Adamax	0.7920	0.613519

In [65]:

```
resultsDf_temp[resultsDf_temp['Log Loss']==resultsDf_temp['Log Loss'].min()]
```

Out[65]:

	Iteration	Method	Accuracy	Log Loss
0	10	RMSProp	0.7925	0.481865

In []:

```
# Predict the Results using 0.5 as a threshold
```

In []:

```
Y_pred = mod.predict_classes(X_test, batch_size=2000, verbose=0)
```

In []:

```
# apply threshold of 0.5
Y_pred = (Y_pred > 0.5)
print(Y_pred)
```

In []:

```
cm1 = confusion_matrix(y_test, Y_pred)
print(cm1)
```

In []:

```
accuracy_model1 = ((cm1[0][0]+cm1[1][1])*100)/(cm1[0][0]+cm1[1][1]+cm1[0][1]+cm1[1][0])
print(accuracy_model1, '% of testing data was classified correctly')
```